

Guneet S. Dhillon

✉ guneet.dhillon@stats.ox.ac.uk
🌐 guneet-dhillon.github.io

Education

- 2021 – 26 **DPhil in Statistics**, *University of Oxford*, UK
- Thesis : Towards Robustness in Predictive and Generative Machine Learning
 - Supervisors : Yee Whye Teh, Tom Rainforth, George Deligiannidis, Arnaud Doucet
- 2014 – 18 **BSc in Computer Science (with Honors) — Turing Scholars Honors & BSc in Mathematics (with Honors)**, *University of Texas at Austin*, TX, USA
- *Certificate in Applied Statistical Modeling*
 - *Minor in Economics*
 - Thesis (Honors) : Training Ising Models on Images using Sparsitron
 - Supervisor : Adam Klivans

Employment

- 2026 **Research Assistant**, *University of Oxford*, Oxford, UK
- 2025 **Research Intern**, *Microsoft Research*, Cambridge, UK
- Publication : E-Scores for (In)Correctness Assessment of Generative Model Outputs
- 2023 – 24 **Machine Learning Intern**, *Boson AI*, Santa Clara, CA, USA
- Publication : L3Ms — Lagrange Large Language Models
- 2019 – 21 **Applied Scientist II**
- 2018 – 19 **& Applied Scientist I**, *Amazon Web Services*, Pasadena, CA, USA
- Publications : Uniform Sampling over Episode Difficulty
A Baseline for Few-Shot Image Classification
 - Patent : Structured Document Analyzer
 - Products : Amazon Rekognition Custom Labels
Amazon Textract
- 2017 **Software Engineer Intern**, *Amazon Web Services*, Palo Alto, CA, USA
- Publication : Stochastic Activation Pruning for Robust Adversarial Defense
- 2016 **Machine Learning Intern**, *CognitiveScale*, Austin, TX, USA
- 2015 **Software Engineer Intern**, *Dell*, Austin, TX, USA

Publications

(* = equal contribution)

- 2026 **E-Scores for (In)Correctness Assessment of Generative Model Outputs**
G. S. Dhillon, J. González, T. Pandeva, A. Curth **AISTATS**
- 2025 **L3Ms — Lagrange Large Language Models**
G. S. Dhillon, X. Shi, Y. W. Teh, A. Smola **ICLR**
- 2024 **On the Expected Size of Conformal Prediction Sets**
G. S. Dhillon, G. Deligiannidis, T. Rainforth **AISTATS**
- 2021 **Uniform Sampling over Episode Difficulty**
S. Arnold, G. S. Dhillon*, A. Ravichandran, S. Soatto* **NeurIPS (Spotlight)**

- 2020 **A Baseline for Few-Shot Image Classification**
G. S. Dhillon, P. Chaudhari, A. Ravichandran, S. Soatto ICLR
- 2018 **Stochastic Activation Pruning for Robust Adversarial Defense**
G. S. Dhillon, K. Azizzadenesheli, Z. Lipton, J. Bernstein, J. Kossaifi, A. Khanna,
A. Anandkumar ICLR

Theses

- 2026 **Towards Robustness in Predictive and Generative Machine Learning**
G. S. Dhillon DPhil
- 2018 **Training Ising Models on Images using Sparsitron**
G. S. Dhillon Bachelor's (Honors)

Patents

- 2020 **Structured Document Analyzer**
G. S. Dhillon, V. Mahadevan, Y. Zhang, M. Wang, G. Payyavula, V. Nguyen, R. Bhotika,
S. Soatto US Patent 10839245

Miscellaneous Articles

(* = equal contribution)

- 2026 **Assessing Generative AI Systems Using E-Scores**
G. S. Dhillon, T. Pandeva, A. Curth ERCIM News 145
- 2020 **Erratum Concerning the Obfuscated Gradients Attack on Stochastic Activation Pruning**
*G. S. Dhillon**, N. Carlini* arXiv

Talks and Presentations

- E-Scores for (In)Correctness Assessment of Generative Model Outputs**
- 2025 Algorithmic Statistics Workshop (University of Oxford)
- 2026 Advances in Adaptive Experimentation Workshop (University College London)
- 2026 Workshop on E-Value Theory, Anytime-Valid Inference, and Game-Theoretic Statistics (University of Twente)
- On the Expected Size of Conformal Prediction Sets**
- 2023 ELLIS Robust Machine Learning Workshop (Aalto University)
- Uniform Sampling over Episode Difficulty**
- 2022 Seminar on Continual Learning, Meta-Learning, and Transfer Learning (Google DeepMind)
- A Baseline for Few-Shot Image Classification**
- 2021 Workshop on Computer Vision with Limited Labels (Amazon CV Conference)
- 2019 Workshop on Meta-Learning (NeurIPS)
- 2019 Workshop on Computer Vision Services/Systems in Amazon (Amazon ML Conference)
- 2019 Workshop on Data-Efficient Learning Techniques for Amazon Scale (Amazon ML Conference)
- Stochastic Activation Pruning for Robust Adversarial Defense**
- 2017 Machine Deception Workshop (NeurIPS)

———— Honors and Awards

University of Oxford

2021 – 26 Clarendon Fund Scholarship *(only 1 awardee from the Department of Statistics)*

University of Texas at Austin

2018 College Scholar, College of Natural Sciences *(only top 20% awardees)*
2017 – 18 Out-of-State Tuition Waiver, College of Natural Sciences *(only 5-7 awardees)*
2017 – 18 Thomas and Elizabeth Merner Scholarship, College of Natural Sciences
2017 – 18 Angus G. and Erna Pearson Endowed Scholarship, Department of Computer Science
2017 College Scholar, College of Natural Sciences *(only top 20% awardees)*
2016 – 17 Motorola Endowed Scholarship, Department of Computer Science
2015 – 16 Angus G. and Erna Pearson Endowed Scholarship, Department of Computer Science
2015 TIDES FRI Summer Research Fellowship, College of Natural Sciences
2014 – 15 Freshman Scholarship, College of Natural Sciences
2014 – 15 Schein Memorial Scholarship, Department of Computer Science

———— Teaching

Module Teaching

2025 Conformal Prediction (University of Oxford)

Teaching Assistant

2022 – 23 Foundations of Statistical Inference (University of Oxford)
2016 Matrices and Matrix Calculations (University of Texas at Austin)

———— Academic Services

Organizing Committee Member

2024 ELLIS Robust Large Language Models Workshop
2024 ELLIS Robust Machine Learning Workshop
2023 ELLIS Robust Machine Learning Workshop

Reviewer

NeurIPS, ICML, ICLR, AISTATS, TMLR